



**BITS Pilani**  
Pilani Campus

# Network Programming

K Hari Babu  
Department of Computer Science & Information Systems



**BITS Pilani**  
Pilani Campus



# History of UNIX & C

# A Brief History of Time (UNIX and C)



- 1969 – First Unix Ken Thompson at AT&T Bell Labs
  - Ideas from Multics:
    - Tree structured file system
    - Program for interpreting commands (shell)
    - Files – unstructured streams of bytes
- 1970 Unix rewritten in assembly for DEC PDP-11
- C – Dennis Ritchie – a systems programming language
  - BCPL → B (Thompson) → C
- 1973 Kernel rewritten in C – eases porting to other machines.
- 1977 Bell Labs released Unix System III
- 1982 AT & T released Unix System V.

# Berkeley Software Division (BSD)



- (1975) Thompson visiting Prof. at UC-Berkeley
- A student Bill Joy added new features
  - Vi editor
  - C shell
  - First paging virtual memory management (Unix) BSD 4.2
  - Sendmail, Pascal compiler
  - Later co-founded Sun Microsystems
- Berkely Software Distribution (BSD)
  - Released BSD 3 in 1979
  - Latest one in 1994 BSD 4.4
- Derivatives
  - Darwin, Free BSD, Net BSD, Open BSD

# Unix after 1979



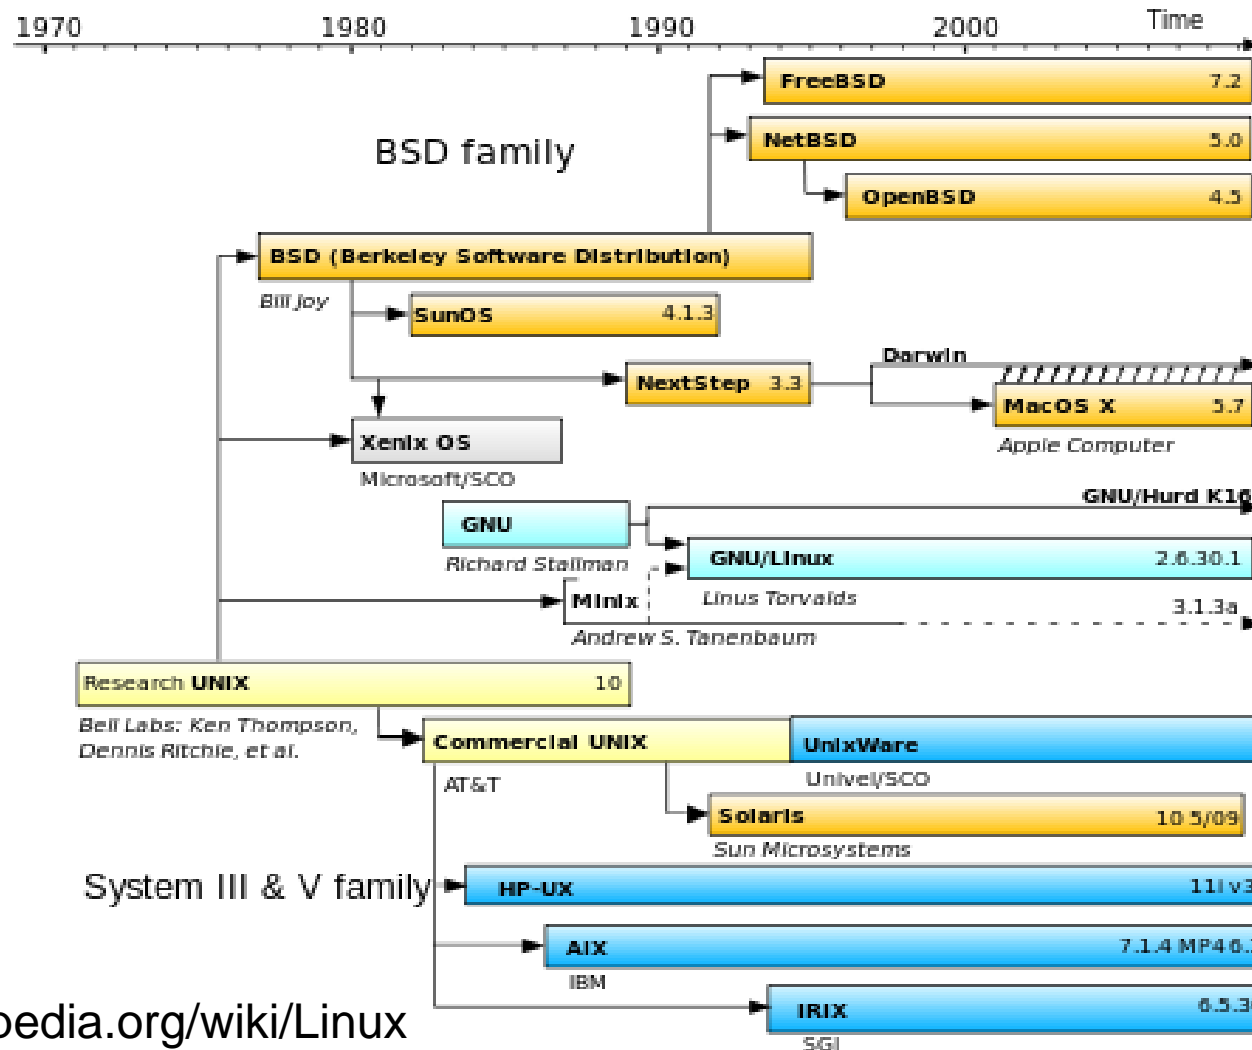
- BSD continued at UC-Berkeley
- Bell Labs System III → Systems V
- POSIX standard (1988)
- Other Software
  - X windows
  - Free Software Foundation
  - GNU Public License
- Minix – (1988)
  - Unix like; MINI-uniX; for education; A. Tannenbaum

# Unix Commercial Versions



- In mid 90s many companies supported high-end features and ported them to their own architecture.
  - Digital's Tru64
  - HP-UX
  - IBM AIX
  - Sequent's DYNIX
  - SGI's IRIX
  - Sun Solaris
  - SunOS

# Unix Timeline



<http://en.wikipedia.org/wiki/Linux>

# Unix Characteristics/Strengths



- Unix is Simple
  - Has only hundreds of systems calls where as other OSs have thousands.
  - Has clear design goals
- Everything is a file
  - Same system calls `open()`, `read()`, `write()` and `close()` can work with files, devices, networks sockets.
- Unix kernel and system utilities are written in C.
  - Easy portability to different architectures.
- Fast process creation time.
- Robust inter process communication (IPC).



- Richard Stallman (1983) Goal a free Unix
  - Known for Free Software movement, GNU, Emacs, gcc
  - Never really released GNU operating system
- Free Software Foundation
  - <http://www.fsf.org/>



<http://en.wikipedia.org/wiki/GNU>

- Minix – (1988)
  - Unix like; MINI-uniX; for education; A. Tannenbaum
- (1991) Linus Torvalds
- For Intel x86 systems
- Moved to big Iron
- more than 90% of today's 500 fastest supercomputers run some variant of Linux
- Network routers
- Embedded systems
- Android


<http://en.wikipedia.org/wiki/Linux>



# Overview of Linux OS Concepts

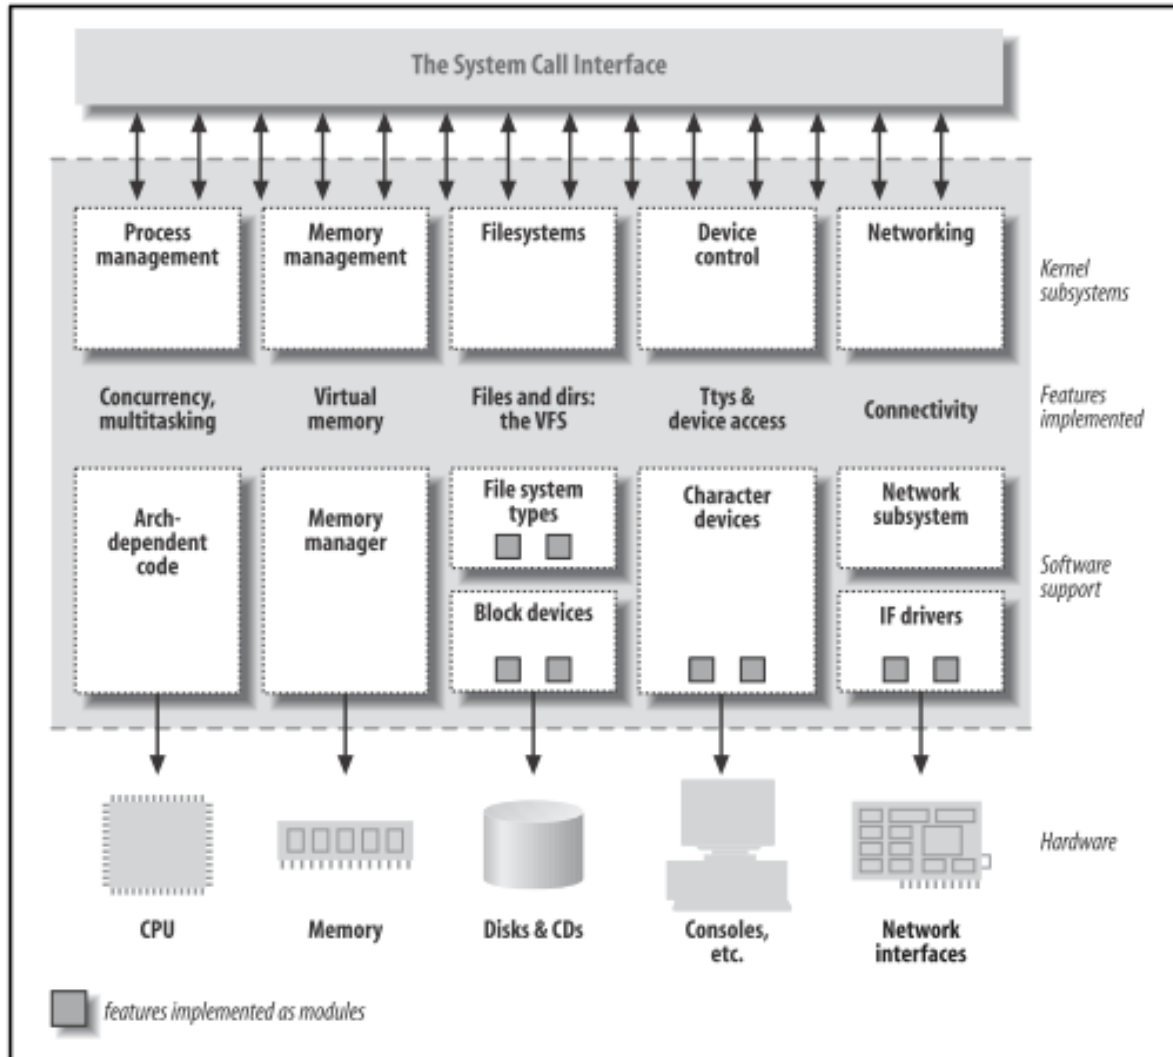
- Kernel
- C library
- Toolchain
- Basic system utilities
- Shell

# About kernel



Source: Wikipedia

# Kernel



Source: Linux Device Drivers

- Central software that manages and allocates computer resources (CPU, RAM, devices etc).
- Kernel greatly simplifies writing programs and using other programs.
  - Processes
  - Virtual address space
  - Virtual CPU

# Kernel Tasks



- Process scheduling
- Memory management
- Provision of a file system
- Creation and termination of processes
- Access to devices
- Networking
- Provision of a system call application programming interface (API)




# Memory Management



- Software size is enormous and RAM is a limited resource.
- Virtual Memory Management
  - Processes are isolated from one another and from the kernel.
  - Only a part of the process needs to be kept in memory. So multiple processes can be held in main memory simultaneously.
- Virtual Address Space
  - On 32 bit systems, there can be 4 GB address space.
  - Normally 2 GB is meant for Kernel and 2 GB for user process.
  - Current process has all the user address space with it.
  - Programmer flexibility.

# Memory Management



Source: wikipedia

- Process Scheduling
  - Preemptive Multitasking
  - Multiple processes can simultaneously reside in memory
  - Preemptive means kernel decides which process and how long it gets CPU not the process itself.
- File System
  - File creation, retrieval, deletion etc.
  - More later
- Networking
  - TCP/IP stack
  - Sending/receiving messages on behalf of user processes.
  - More later

# Kernel Tasks: Process Management




- Kernel can load new program into memory, provide it with the resources.
  - `fork()`
  - `execve()`
- Instance of a running program is termed as a process.
- Processes need to invoke system calls to access resources.
- Kernel executes the system calls on behalf of the process in kernel space.
- **Access to devices:**
  - Kernel provides an interface that standardizes and simplifies access to devices.
  - Keyboard, display etc.

# Kernel Tasks: System Calls



- Kernel entry points are known as system calls.
- System calls are software interrupts.

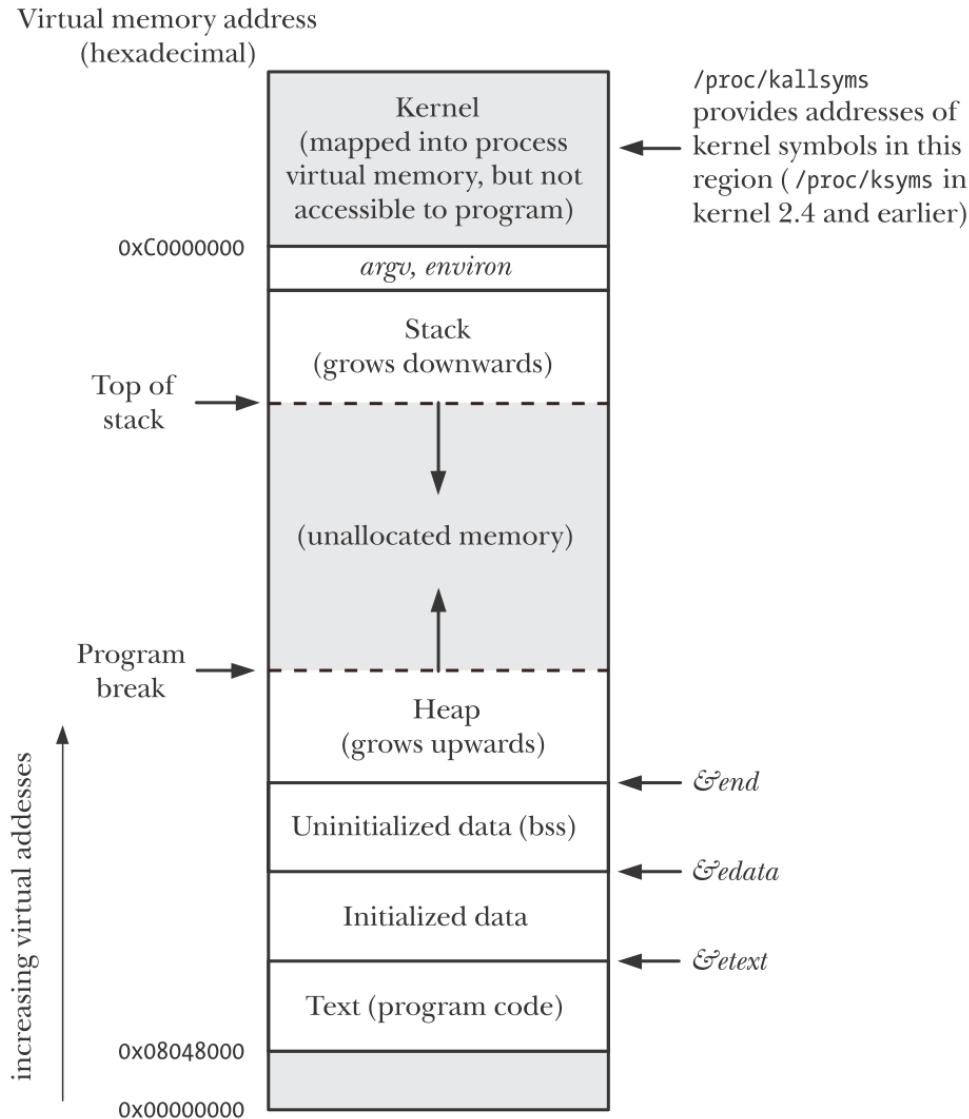


# Kernel Mode vs User Mode



- Processor architecture allows the CPU to operate in at least two different modes:
  - User mode
  - Kernel mode
- Similarly virtual memory is marked as user space and kernel space.
  - When running user mode, the CPU can access only user space.
  - When running in kernel mode, the CPU can access both user and kernel space.
- By this arrangement, user processes can't access instructions and data structures of kernel.

# Kernel Space vs User Space





**BITS Pilani**  
Pilani Campus



# Sys Calls & C Library

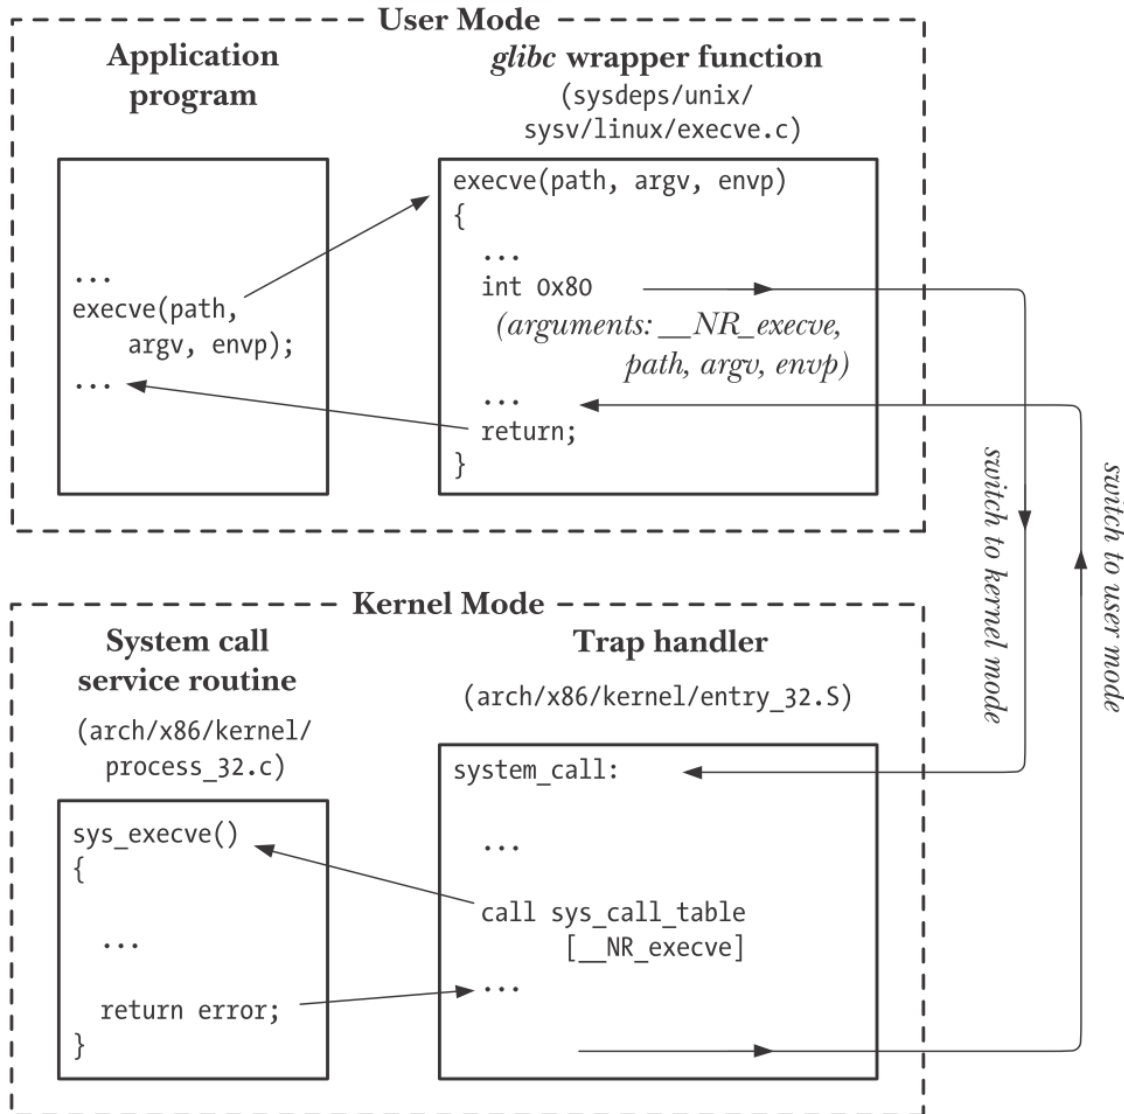


# Kernel Tasks: System Calls



- A system call changes the processor from user mode to kernel mode.
- The set of sys calls are fixed. Each system call is identified by a unique number.
- System calls arguments are copied from user space to kernel space (registers).
- Wrapper function executes a trap instruction (int 0x80) causing the CPU to enter kernel mode.
- Example: `execve()`
  - System call number is 11.
  - `Sys_call_table` contains the address of the routine `sys_execve()`.
- Not all C library calls invoke system calls. E.g.: `<string.h>`

# Kernel Tasks: System Calls



# Sys Calls & C Library



- System call is a controlled entry point into the kernel.
- For every sys call there is a wrapper function in C library.
  - All library functions are not sys calls.
  - We use wrapper functions in programs
  - `fork()`, `execve()` ...
- Wrapper function copies the arguments into specific registers. Also copies sys call number into a specific CPU register.
- Wrapper functions executes trap instruction `int 0x80`.
  - Causes CPU to switch from user mode to kernel mode.
- Kernel executes *system\_call* routine at *arch/x86/kernel/entry.S*

# Sys Calls & C Library



- On i386, the parameters of a system call are transported via registers.
  - The system call number goes into %eax
  - the first parameter in %ebx
  - the second in %ecx
  - the third in %edx
  - the fourth in %esi
  - the fifth in %edi
  - the sixth in %ebp.

# Sys Calls & C Library



- *system\_call()* routine:
  - Saves register values onto the stack
  - Checks the validity of the system call number.
  - Invokes the corresponding service routine.
  - Service routine returns a result status to the *system\_call* routine.
  - Restores register values from the kernel stack and places the system call return value on the user process stack.
  - Returns to the wrapper function, simultaneously returning the process to enter user mode.
- Wrapper function checks the return value and if it is an error it sets the *errno* variable.
- System calls incur an appreciable overhead.
  - Calling a C library wrapper function is synonymous with invoking the corresponding system call routine.

- Many library functions do not make use of system calls.
- Often library functions provide a more caller-friendly interface than the underlying sys call.
  - `fopen()` uses `open()`
  - `printf` uses `write()`
  - `malloc()` uses `brk()`
- GNU C library (*glibc*)
  - `libc.so.6`
  - Where is it stored? List dynamic dependencies
    - `ldd a.out`
  - Finding the version
    - `./libc.so.6`



**BITS Pilani**  
Pilani Campus



# Handling Errors

# Errors from System Calls



- A service routine in kernel returns a negative number in case of error. The negative number corresponds to standard error codes.

`/usr/include/asm-generic/errno-base.h`

```
#ifndef _ASM_GENERIC_ERRNO_BASE_H
#define _ASM_GENERIC_ERRNO_BASE_H

#define EPERM          1  /* Operation not permitted */
#define ENOENT          2  /* No such file or directory */
#define ESRCH          3  /* No such process */
#define EINTR          4  /* Interrupted system call */
#define EIO             5  /* I/O error */
#define ENXIO           6  /* No such device or address */
#define E2BIG           7  /* Argument list too long */
#define ENOEXEC         8  /* Exec format error */
#define EBADF           9  /* Bad file number */
#define ECHILD          10 /* No child processes */
#define EAGAIN          11 /* Try again */
#define ENOMEM          12 /* Out of memory */
#define EACCES          13 /* Permission denied */
#define EFAULT          14 /* Bad address */
#define ENOTBLK         15 /* Block device required */
#define EBUSY           16 /* Device or resource busy */
```



# Errors from System Calls



- In case of error, wrapper function stores the positive value of the error code into `errno` variable and returns -1.

```
2 cnt = read(fd, buf, numbytes);
3 if (cnt == -1) {
4     if (errno == EINTR)
5         fprintf(stderr, "read was interrupted by a signal\n");
6     else {
7         /* Some other error occurred */
8     }
9 }
10
```

- `perror()` and `strerror()` can be used to print the error.

```
2 fd = open(pathname, flags, mode);
3 if (fd == -1) {
4     perror("open");
5     exit(EXIT_FAILURE);
6 }
7
```

# errno variable



- It is present one per each process.
- It is set in wrapper function after sys call returns error.
  - Every time it is over written.
  - So only after a sys call returns -1, we refer to errno.

# Handling Errors from Library Functions



- Some library functions return error information exactly like system calls.
  - Return -1
  - errno is set
- Some return value other than -1 but set errno
  - fopen
- Others do not use errno at all
  - gethostbyname

# Tracing System Calls



- *strace* command allows us to trace the system calls made by a program.

```
haribabuk@haribabuk-VirtualBox ~ $ strace ./a.out
execve("./a.out", ["/a.out"], [/* 48 vars */]) = 0
brk(0)                                = 0x10e9000
access("/etc/ld.so.nohwcap", F_OK)    = -1 ENOENT (No such file or directory)
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7fd26b5e2000
access("/etc/ld.so.preload", R_OK)    = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY)    = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=128950, ...}) = 0
mmap(NULL, 128950, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7fd26b5c2000
close(3)                               = 0
```

- Each system call is displayed with input and output arguments
- Arguments are printed in symbolic form.
- *ltrace* is used for tracing library calls.

```
haribabuk@haribabuk-VirtualBox ~ $ ltrace ./a.out
__libc_start_main(0x4005f4, 1, 0x7fffd10e0ba8, 0x400700, 0x400790 <unfinished ...>
socket(2, 3, 1)                                = -1
recvfrom(0xffffffff, 0x7fffd10e04d0, 1500, 0, 0) = -1
+++ exited (status 0) +++
```

# Q&A





**BITS Pilani**  
Pilani Campus



**Thank You**