



**BITS Pilani**  
Pilani Campus

# Network Programming

K Hari Babu  
Department of Computer Science & Information Systems



**BITS Pilani**  
Pilani Campus



# File I/O

# File Descriptor (fd)



- File Descriptors refer to all types of open files.
  - Pipes, FIFOs, sockets, terminals, devices, and regular files.
- Each process has its own set of file descriptors.
- All system calls refer to file descriptors for performing I/O.
- Three standard file descriptors.

File descriptor	Purpose	POSIX name	<i>stdio</i> stream
0	standard input	STDIN_FILENO	<i>stdin</i>
1	standard output	STDOUT_FILENO	<i>stdout</i>
2	standard error	STDERR_FILENO	<i>stderr</i>

- These three descriptors are open in the shell process.
  - Whenever a new program is executed in the shell, a child process is created. All three descriptors remain open in the child.
- File descriptors are different from FILE streams. FILE stream is a C library abstraction over fd.

# File I/O: open()



```
2 #include <sys/stat.h>
3 #include <fcntl.h>
4 int open(const char * pathname , int flags , ... /* mode_t mode */);
5 Returns file descriptor on success, or -1 on error
```

- *flags* is a bitmask that refers to read only or write only or both.
- *mode* refers to permissions. *mode* is used only when creating a file.

```
2 openFlags = O_CREAT | O_WRONLY | O_TRUNC;
3 filePerms = S_IRUSR | S_IWUSR | S_IRGRP | S_IWGRP |
4             S_IROTH | S_IWOTH;          /* rw-rw-rw- */
5 outputFd = open(argv[2], openFlags, filePerms);
6 if (outputFd == -1)
7     errExit("opening file %s", argv[2]);
```

- O\_CREAT option is used when a new file is to be created.
- O\_TRUNC option is used when the data in the file has to be deleted.
- O\_APPEND is used for appending to the existing file.
- Several other flags also present ... (R1: 4.3.1)

# File I/O: read()



```
2  #include <unistd.h>
3  ssize_t read(int  fd , void * buffer , size_t  count );
4  /*Returns number of bytes read, 0 on EOF, or -1 on error*/
```

- Reads at most *count* bytes from the open file referred to by *fd* and stores them in a *buffer*.
- It returns the no of bytes actually read or EOF or -1.
  - *count*: maximum number of bytes to read
  - *buffer*: address of the memory buffer into which the input data is to be placed.
  - Read may read less than *count*.
    - In regular files, we may be close to EOF.
    - In pipes, FIFOs, and sockets it may read less than *count* due to non-availability of data.

# File I/O: read()



```
2 char buffer[MAX_READ + 1];
3 ssize_t numRead;
4 numRead = read(STDIN_FILENO, buffer, MAX_READ);
5 if (numRead == -1)
6     errExit("read");
7 buffer[numRead] = '\0';
8 printf("The input data was: %s\n", buffer);
```

- STDIN\_FILENO refers to fd 0.
- At line 9, we need to include NULL character because *read()* doesn't do it itself.

# File I/O: write()



```
2  #include <unistd.h>
3  ssize_t write(int fd, void * buffer , size_t  count );
4  /*Returns number of bytes written, or -1 on error */
```

- Writes up to *count* bytes from *buffer* to the open file referred to by *fd*.
- Returns the number of bytes actually written which may be less than *count*.

# File I/O: close()



```
2  #include <unistd.h>
3  int close(int  fd );
4  /*Returns 0 on success, or -1 on error*/
```

- It is called after all I/O has been completed in order to release the file descriptor *fd* and its associated kernel resources.
- When a process terminates all of its open file descriptors are automatically closed.



# File I/O: lseek()



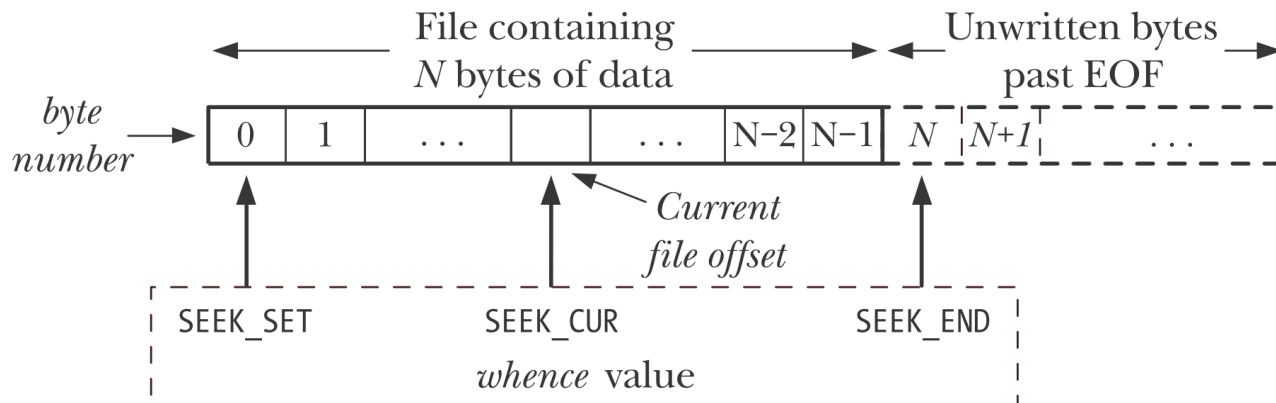
```
2  #include <unistd.h>
3  off_t lseek(int fd , off_t offset , int whence );
4  /*Returns new file offset if successful, or -1 on error*/
```

- It adjusts the file offset of the open file referred to by the *fd*, according to the values specified in *offset* and *whence*.
- Kernel records a *file offset* for each open file.
- This is the location in the file at which the next read or write will commence.
- When the file opened the offset is set to 0 i.e. the beginning of the file.

# File I/O: lseek()



- *whence* argument can be
  - SEEK\_SET
  - SEEK\_CUR
  - SEEK\_END



- `lseek()` simply adjusts the file offset, it doesn't cause any physical device access.
- We can't apply `lseek()` to pipe, FIFO, socket or terminal.

# File I/O: lseek()



```
2 curr = lseek(fd, 0, SEEK_CUR); /* Retrives the file offset*/
3 lseek(fd, 0, SEEK_SET);        /* Start of file */
4 lseek(fd, 0, SEEK_END);        /* Next byte after the end of the file */
5 lseek(fd, -1, SEEK_END);       /* Last byte of file */
6 lseek(fd, -10, SEEK_CUR);      /* Ten bytes prior to current location */
7 lseek(fd, 10000, SEEK_END);    /* 10001 bytes past last byte of file */
```

- File holes
  - What if we read after `lseek(fd, 10000, SEEK_END)`?
    - Returns 0.
  - What if we write after `lseek(fd, 10000, SEEK_END)`?
    - It creates a file hole. File holes do not take disk space.
- File holes are useful when a program need to access a wide range of addresses (offset) but is unlikely to touch all of the potential blocks.
  - Virtual hard disks

# Universality of I/O



- The same four system calls – *open()*, *read()*, *write()*, and *close()* – are used to perform I/O on all types of files.
  - Regular files, Pipe, FIFO, sockets, terminal devices
- *ioctl()* system call is for operations that fall outside the universal I/O model.

```
2  #include <sys/ioctl.h>
3  int ioctl(int fd , int request , ... /* argp */);
4  /*Value returned on success depends on request, or -1 on error*/
```

- *fd* refers to any file or device.
- *request* refers to the constant specific to the device.
- *argp* is the value or buffer depending the type of request.

# Universality of I/O: *ioctl()*

- e.g. for updating flags of inode, *ioctl()* is used.

```
2 int attr;
3 if (ioctl(fd, FS_IOC_GETFLAGS, &attr) == -1)    /* Fetch current flags */
4     errExit("ioctl");
5 attr |= FS_NOATIME_FL; /*do not update last file access time*/
6 if (ioctl(fd, FS_IOC_SETFLAGS, &attr) == -1)    /* Update flags */
7     errExit("ioctl");
```

# File Descriptors and Open Files



- It is possible to have multiple descriptors referring to the same open file.
- There are three data structures maintained by the kernel:
  - The per-process file descriptor
  - The system-wide table of open file descriptions
  - The file system i-node table.
- Per-process file descriptor table
  - Flags (close-on-exec)
  - A reference to the open file description

**innovate**      **achieve**      **lead**



# File Descriptors and Open Files



- System-wide table of all open file descriptions
  - The current file offset
    - Modified by `read()`, `write()` and `lseek()`
  - Status flags (flags argument to `open()`)
  - File access mode (read-only, write only etc as specified in `open()`)
  - Settings related to signal driven I/O
  - a reference to i-node object for this file.
- i-node table
  - Each file system has a table of i-nodes for all files residing in the file system.
  - File type (regular file, socket, FIFO etc)
  - A pointer to list of blocks
  - Various properties of the file (size, timestamps etc)



# File Descriptors and Open Files



- Two descriptors in different process may refer to the same open file entry.
  - `fork()`
  - Passing descriptor using UNIX domain sockets
- Two open file entries can refer to same i-node.
  - When the same is open twice in the same process or in different processes.
- When an open file entry is shared
  - Updating file offset or flags effects the other process.
- `close-on-exec` flag individual to a fd. Changing doesn't effect the other processes.

# Duplicating File Descriptors: dup()



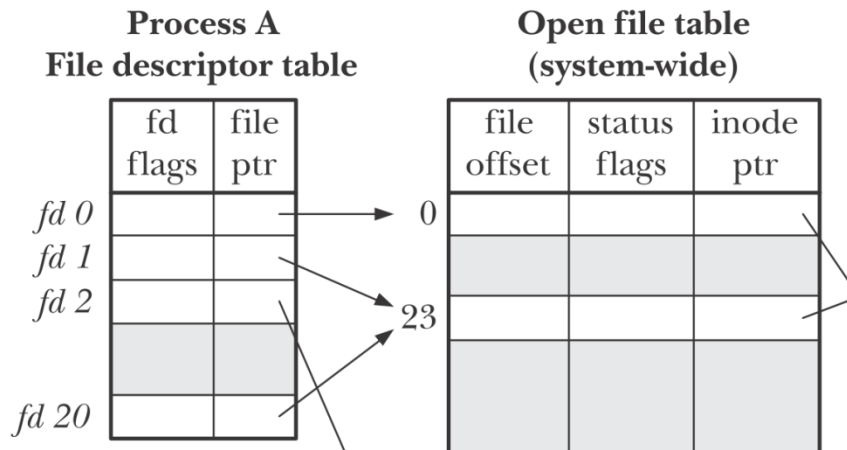
```
$ ./myscript > results.txt 2>&1
```

- Here 2>&1 indicates that standard error (fd 2) to be sent to the same place where standard output (fd 1) is being sent.
- This is possible by duplicating fd 2 to refer to open table entry referred by fd 1.

```
2  #include <unistd.h>
3  int dup(int oldfd );
4  /*Returns (new) file descriptor on success, or -1 on error*/
```

- Dup takes *oldfd* and returns a new fd that refers to the same open file table entry.
- New fd is guaranteed to be the lowest unused file descriptor.

# Duplicating File Descriptors: dup2()



- `close(1); dup(20)`

```
2  #include <unistd.h>
3  int dup2(int  oldfd , int  newfd );
4  /*Returns (new) file descriptor on success, or -1 on error*/
```

- If *newfd* is open it closes it and copies the pointer in *oldfd* to *newfd* slot.
- This is done atomically.

# File Control Operations: *fcntl()*



- `fcntl()` call performs operations on open file descriptors.

```
2  #include <fcntl.h>
3  int fcntl(int fd , int cmd , ...);
4  /*Return on success depends on cmd, or -1 on error*/
```

- `cmd` refers to commands.
- e.g. to change the flag after opening file

```
2  int flags;
3  flags = fcntl(fd, F_GETFL);
4  if (flags == -1)
5      errExit("fcntl");
6  flags |= O_APPEND;
7  if (fcntl(fd, F_SETFL, flags) == -1)
8      errExit("fcntl");
```

- Append is flag is being added to the flags in open file table entry.



# I/O Buffering

- For speed and efficiency, I/O systems calls and I/O library calls buffer data.
- Two levels of buffering
  - Kernel buffer cache
    - Makes read and write calls faster
    - Reduces number of disk access by kernel
  - Buffering in the standard i/o library (optional)
    - Reduces number of system calls to access data

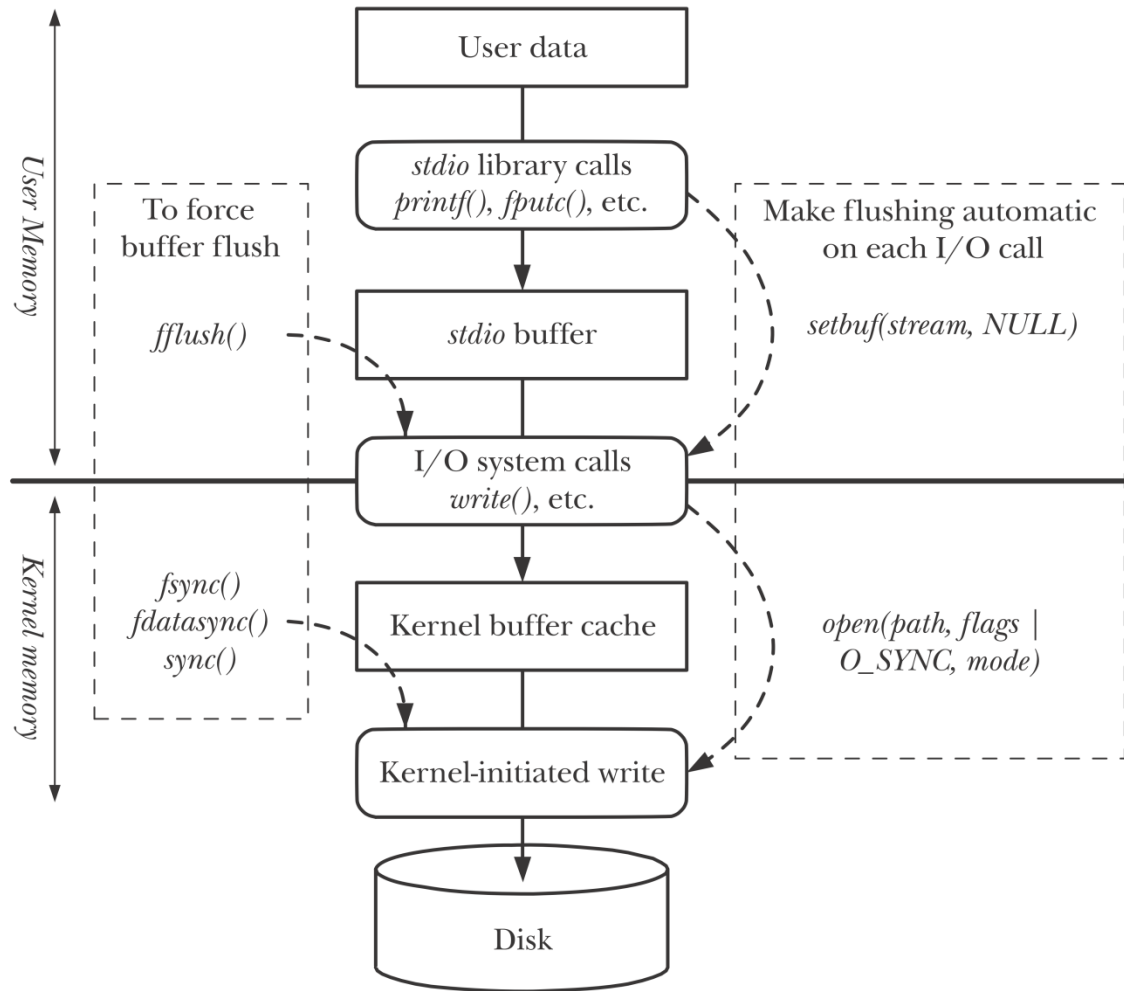
```
2 write(fd, "abc", 3);
```

- This call can't directly write to the disk. This writes to a buffer in the kernel. Kernel later syncs these contents with the disk.

```
2 read(fd, buf, 3);
```

- This call transfers 3 bytes of data from kernel buffer to user buffer *buf*.

# Writing Data to a File



# Kernel Buffering of File I/O



- Buffer cache:
  - Set of buffers Kernel maintains to store disk blocks. Seize of buffer cache is adapted as per the availability of the physical memory.
  - When a read() call is issued, Kernel reads the disk block from the disk and stores it in a buffer.
  - Data is copied from buffer cache to the buffer in the user space.
  - Similarly when a user process writes, kernel writes to the buffer.
  - Kernel periodically syncs dirty buffers with disk.
- This allows read() and write to be faster.



# Buffering in *stdio* Library



- C library buffers the data to reduce the number system calls (read, write). *fopen()* call opens a buffered stream for a file.

```
2  #include <stdio.h>
3  int setvbuf(FILE * stream , char * buf , int mode , size_t size );
4  /*Returns 0 on success, or nonzero on error*/
```

- This is a library function that controls the type of buffering.
- This function must be called before any I/O operation.
- If *buf* is null, stdio automatically allocates the buffer for use with the *stream*.
- mode
  - \_IONBF: no buffering. E.g. stderr
  - \_IOLBF: line buffering. Default for terminal devices. Output is buffered until newline char. Data is read a line at a time.
  - \_IOFBF: fully buffered I/O. data is read or written in units of buffer size. Default for disk files.

# Flushing a stdio Buffer



```
2  #include <stdio.h>
3  int fflush(FILE * stream );
4  /*Returns 0 on success, EOF on error*/
```

- Regardless of the current buffering mode, at any time, we can force the data to written.
- fflush() function flushes the output buffer for that particular FILE stream.
  - It can be used on input stream also.

# Controlling Kernel Buffering of File I/O



- Two type of synchronization
  - Synchronized I/O file integrity
    - Both data and meta data about the file are synchronized with disk.
  - Synchronized I/O data integrity
    - Only data is synchronized with the disk.

```
2  #include <unistd.h>
3  int fsync(int  fd );
4  /*Returns 0 on success, or -1 on error*/
```

- Flushes both data and metadata such as file size, time stamps etc associated with *fd*.

```
2  #include <unistd.h>
3  int fdatasync(int  fd );
4  /*Returns 0 on success, or -1 on error*/
```

- Flushes only data buffers of file descriptor *fd*.

# Controlling Kernel Buffering of File I/O



```
2 fd = open(pathname, O_WRONLY | O_SYNC);
```

- If we use O\_SYNC flag while opening a file, after every write, both data and metadata will be flushed to disk.

```
2 #include <unistd.h>
3 void sync(void);
```

- It causes all kernel buffers containing modified data including metadata to be flushed to disk.
- Call returns only after syncing.

# Mixing Library Functions and I/O Sys Calls



- Sys calls: `open()`, `read()`, `write()`, `close()`
  - Work with file descriptors
- Library functions: `fopen()`, `fprintf()`, `fscanf()`, `fclose()`, ...
  - Work with FILE streams

```
2  #include <stdio.h>
3  int fileno(FILE * stream );
4  /*Returns file descriptor on success, or -1 on error*/
5  FILE *fdopen(int fd , const char * mode );
6  /*Returns (new) file pointer on success, or NULL on error*/
```

- Given a stream, `fileno()` returns the corresponding `fd`
  - Given a file descriptor, `fdopen()` creates corresponding FILE stream.
- `fdopen()` is useful while dealing with pipes and sockets.

# Q&A





**BITS Pilani**  
Pilani Campus



**Thank You**