



BITS Pilani
Pilani Campus

Network Programming

K Hari Babu
Department of Computer Science & Information Systems



Signals (R1: Ch20)

What is a Signal?



- A signal is a notification to a process that an event has occurred.
 - A signal is an *asynchronous* event which is delivered to a process.
 - Asynchronous means that the event can occur at any time
 - e.g. user types `ctrl-C`
- Source of signals:
 - Hardware exceptions
 - Dividing by 0, accessing inaccessible memory
 - User typing special characters at the terminal
 - Control-c, Control-\, Control-Z
 - Software events
 - Data available on a descriptor
 - A timer went off
 - Child is terminated etc.
 - kill function allows a process to send a signal to another process.

- Each signal is defined by a unique integer, starting from 1 to 31. 0 is a NULL signal.
 - Mapping varies across architectures. Better go by symbols.

Name	Signal number	Description	SUSv3	Default
SIGABRT	6	Abort process	•	core
SIGALRM	14	Real-time timer expired	•	term
SIGBUS	7 (SAMP=10)	Memory access error	•	core
SIGCHLD	17 (SA=20, MP=18)	Child terminated or stopped	•	ignore
SIGCONT	18 (SA=19, M=25, P=26)	Continue if stopped	•	cont
SIGEMT	undef (SAMP=7)	Hardware fault		term
SIGFPE	8	Arithmetic exception	•	core
SIGHUP	1	Hangup	•	term
SIGILL	4	Illegal instruction	•	core
SIGINT	2	Terminal interrupt	•	term
SIGIO / SIGPOLL	29 (SA=23, MP=22)	I/O possible	•	term
SIGKILL	9	Sure kill	•	term
SIGPIPE	13	Broken pipe	•	term

SAMP: Sun SPARC and SPARC64 (S), HP/Compaq/Digital Alpha (A), MIPS (M), and HP PA-RISC (P)

Important Signals



- **SIGABRT**
 - When a process calls *abort()*, this signal is delivered to the process. Terminates with a core dump.
- **SIGALRM**
 - Kernel generates this signal upon the expiration of a timer set by *alarm()* or *settimer()*.
- **SIGCHLD**
 - Kernel generates this signal upon the termination of a child process.
- **SIGHUP**
 - When a terminal is disconnected, the controlling process of the terminal is sent this signal.
 - Daemons process repond to this signal by reinitializing from config file.

Important Signals



- **SIGINT**
 - Generated when a user presses Ctrl-c on a terminal to interrupt a process.
- **SIGIO**
 - Useful in signal driven I/O on sockets.
- **SIGKILL**
 - Can't be blocked, ignored or caught by a handler. Always terminates a process. Used by admins.
- **SIGPIPE**
 - Kernel generates this when a pipe has no readers but a process writes into it.
- **SIGQUIT**
 - Generated when user presses Ctrl-\ on the terminal. It terminates the process with core dump.

Important Signals



- **SIGSEGV**
 - Generated when
 - Referencing an unmapped address
 - Updating a read-only page.
 - Accessing kernel memory.
- **SIGSTOP**
 - Used by admin to stop a process. Can't be ignored, blocked or handled. Process can be started again by SIGCONT signal.
- **SIGTERM**
 - Standard signal used for terminating a process. *init* process sends this signal during shutdown.
- **SIGUSR1 & SIGUSR2**
 - Kernel never generates these signals.
 - Available for programmer defined purposes.

- A signal is said to be *generated* by some event. Once generated signal is *delivered* to the process. Between the time it is generated and delivered, signal is said to be *pending*.
 - A pending signal is delivered as soon as the process scheduled to run or immediately if the process is running.
- Sometimes we do not want signals to be delivered to the process when executing a critical code segment.
 - Signals can be added to *signal mask* in the kernel. These signals are blocked. If a such is signal is generated, it will be kept pending.
 - When the mask is cleared, the pending signals are delivered to the process.

Signal Disposition



- A process can inform kernel how it wants to deal with a signal:
 - ignore/discard the signal (not possible with `SIGKILL` or `SIGSTOP`)
 - Kernel will not deliver such a signal to the process.
 - Catch the signal and execute a signal handler function, and then possibly resume execution.
 - Process installs a handler function for a signal with the Kernel. When the signal is received, kernel executes the function on the process behalf in user space.
 - Let the default action apply. Every signal has a default action.
 - Default action can be ignore, terminate the process. Depends on the signal.
- This choice is called the *signal disposition*.

Setting Signal Disposition



- *signal()* system call takes a function pointer *handler* and registers against the signal *signo*.

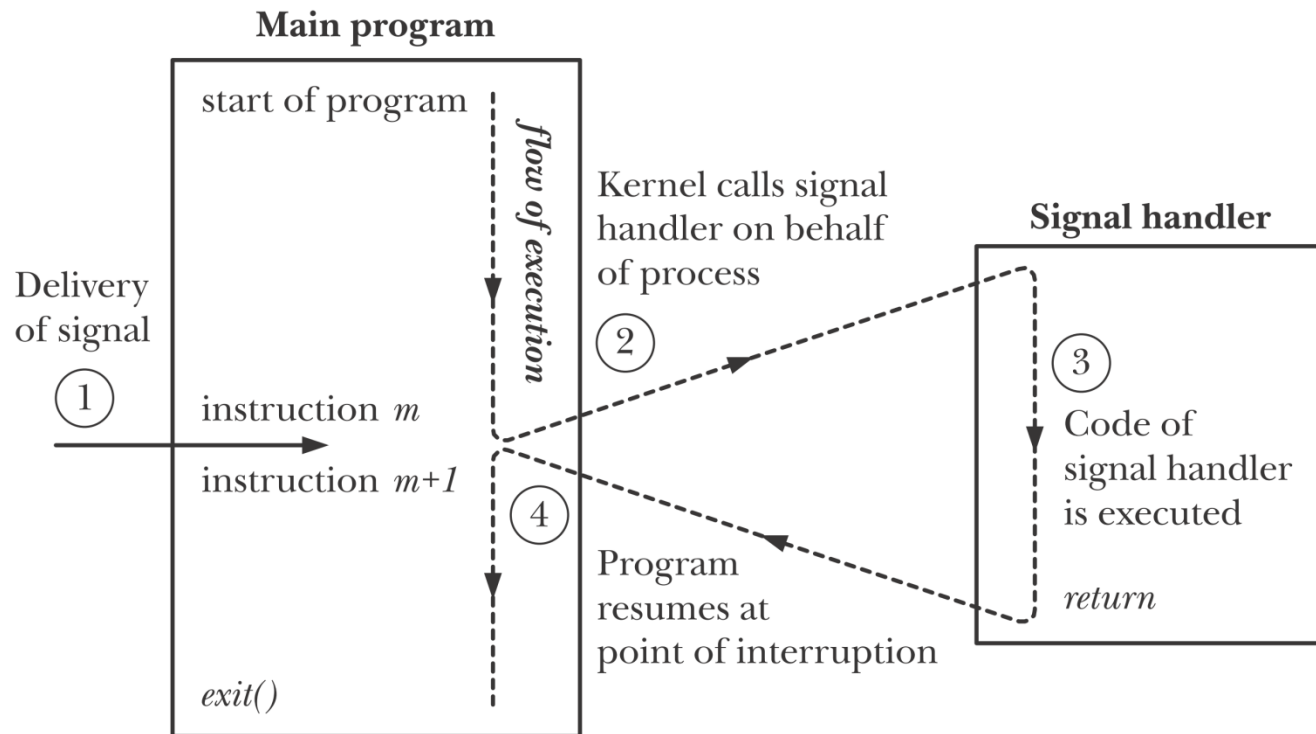
```
2  #include <signal.h>
3  typedef void Sigfunc(int); /* my defn */
4  Sigfunc *signal( int signo, Sigfunc *handler );
5  /*Returns previous signal disposition if ok, SIG_ERR on error.*/
```

- For other dispositions, the Sigfunc values are
 - SIG_IGN Ignore / discard the signal.
 - SIG_DFL Use default action to handle signal.
- In case of error
 - SIG_ERR Returned by signal() as an error.
- We can't know the current disposition for a signal with this sys call.
 - *sigaction()* sys call can do that.

Signal Handlers



- Kernel calls the handlers on the process's behalf.
- Handler may be called at any time.
 - After the execution of the handler, program resumes from the point where it got interrupted.



Signal Handler: example



```
2  ▾ /*signals/ouch.c*/
3  #include <signal.h>
4  #include "tlpi_hdr.h"
5  static void
6  sigHandler(int sig)
7  ▾ {
8      printf("Ouch!\n");          /* UNSAFE (see Section R1: 21.1.2) */
9  }
10 int
11 main(int argc, char *argv[])
```

```
1  $ ./ouch
2  0                               Main program loops, displaying successive integers
3  Type Control-C
4  Ouch!                           Signal handler is executed, and returns
5  1                               Control has returned to main program
6  2
7  Type Control-C again
8  Ouch!
9  3
10 Type Control-\ (the terminal quit character)
11 Quit (core dumped)
```

kill() and raise()function



- Send a signal to a process (or group of processes).

```
2  #include <signal.h>
3  int kill( pid_t pid, int signo );
4  int raise(int signo);
5  /*Return 0 if successful, -1 on error.*/
```

- pid > 0 send signal to process pid
 - pid== 0 send signal to all processes whose process group ID equals the sender's pgid.
 - e.g. parent kills all children
- Using raise(), a process can send a signal to itself.
- To know whether a PID is in use
 - Send a null signal to that PID
 - kill(PID, 0)
- kill command
 - kill -INT 9400

Signal Sets



- Many signal related system calls need a set of signals as input.
- Signal set is a data structure represented by *sigset_t* data type.
- Following are the library functions on *sigset_t* data type.

```
2  #include <signal.h>
3  int sigemptyset(sigset_t *set);
4  int sigfillset(sigset_t *set);
5  int sigaddset(sigset_t *set, int signo);
6  int sigdelset(sigset_t *set, int signo);
7  /*All four return: 0 on success, -1 on error */
8  int sigismember(const sigset_t *set, int signo);
9  /*Returns: 1 if true, 0 if false, -1 on error*/
```

- For each process kernel maintains a *signal mask* – set of signals the process has blocked.
 - When a signal is to be delivered but it is blocked, then that signal is kept pending until it is unblocked by the process.
 - When a signal handler is invoked, the signal that caused invocation is automatically added to the mask.
- Using *sigprocmask()* system call, signals can be added or deleted from the mask or retrieve the mask.

```
2  #include <signal.h>
3  int sigprocmask(int how , const sigset_t * set , sigset_t * oldset );
4  /*Returns 0 on success, or -1 on error*/
```

- how is interpreted as
 - SIG_BLOCK: add *set* to the signal mask
 - SIG_UNBLOCK: delete *set* from the mask.
 - SIG_SETMASK: reset signal mask to *set*.

sigprocmask()



```
2  sigset_t blockSet, prevMask;
3  /* Initialize a signal set to contain SIGINT */
4  sigemptyset(&blockSet);
5  sigaddset(&blockSet, SIGINT);
6  /* Block SIGINT, save previous signal mask */
7  if (sigprocmask(SIG_BLOCK, &blockSet, &prevMask) == -1)
8      errExit("sigprocmask1");
9  /* ... Code that should not be interrupted by SIGINT ... */
10 /* Restore previous signal mask, unblocking SIGINT */
11 if (sigprocmask(SIG_SETMASK, &prevMask, NULL) == -1)
12     errExit("sigprocmask2");
```


Pending Signals



- To know pending signals we use *sigpending()* system call.

```
2  #include <signal.h>
3  int sigpending(sigset_t * set );
4  /*Returns 0 on success, or -1 on error*/
```

- Pending signals are returned through *set*.
 - They are examined using *sigismember()* function.
- If a blocked signal is generated more than once then in most systems the signal is delivered only once. That is the signal is not queued.
- If many signals of *different* types are ready to be delivered (e.g. a `SIGINT`, `SIGSEGV`, `SIGUSR1`), they are not delivered in any fixed order.

Setting Signal Disposition: *sigaction()*



```
2  #include <signal.h>
3  int sigaction(int sig , const struct sigaction * act ,
4  struct sigaction * oldact );
5  /*Returns 0 on success, or -1 on error*/
6
7  struct sigaction {
8      void (*sa_handler)(int);/* Address of handler */
9      sigset_t sa_mask;        /* Signals blocked during handler
10                               invocation */
11      int sa_flags;            /* Flags controlling handler invocation */
12      void (*sa_restorer)(void);/* Not for application use */
13  };
```

- An alternative to *signal()*
 - Allows us to retrieve the disposition of a signal without changing it.
 - To atomically block signals while executing in a handler.
 - To retrieve attribute of a signal such pid, uid etc using SA_SIGINFO.
 - To automatically restart a system call using SA_RESTART.

Waiting for a Signal



- Calling *pause()* suspends execution of the process until the call is interrupted by a signal handler.

```
2  #include <unistd.h>
3  int pause(void);
4  /*Always returns -1 with errno set to  EINTR*/
```

- When a signal is handled, *pause()* is interrupted and always returns -1.
- *sigsuspend()*: atomically unblocks signals and suspends execution of the process until a signal is caught and its handler returns.

```
2  #include <signal.h>
3  int sigsuspend(const sigset_t * mask );
4  /*(Normally) returns -1 with errno set to  EINTR*/
```

- Replaces the signal mask in the kernel with *mask*.
 - Suspends execution until signal handler returns.

pause() example



```
2  #include <signal.h>
3  long n;
4  void sigalrm(int signo){
5      alarm(1);
6      n=n+1;
7      printf("%d seconds elapsed\n",n);
8  }
9  main(){
10     signal(SIGALRM, sigalrm);
11     alarm(1);
12
13     while(1) pause();
14 }
```

Synchronously Waiting for a Signal



- *pause()* and *sigsuspend()* wait until a signal handler is executed. But we can do away with signal handlers.

```
2  #define _POSIX_C_SOURCE 199309
3  #include <signal.h>
4  int sigwaitinfo(const sigset_t * set , siginfo_t * info );
5  /*Returns number of delivered signal on success, or -1 on error*/
```

- Suspends execution until one of the signals in *set* becomes pending and returns that signal number.
- Useful only if signal is *set* are blocked using *sigprocmask()*.
- Faster than *sigsuspend()* because there is no signal handler.

sigwaitinfo() example

innovate

achieve

lead

```
1 ▾ /*signals/t_sigwaitinfo.c*/
2 ▾ /* Block all signals (except SIGKILL and SIGSTOP) */
3     sigfillset(&allSigs);
4     if (sigprocmask(SIG_SETMASK, &allSigs, NULL) == -1)
5         errExit("sigprocmask");
6     printf("%s: signals blocked\n", argv[0]);
7
8 ▾     for (;;) { /* Fetch signals until SIGINT (^C) or SIGTERM */
9         sig = sigwaitinfo(&allSigs, &si);
10        if (sig == -1)
11            errExit("sigwaitinfo");
12        if (sig == SIGINT || sig == SIGTERM)
13            exit(EXIT_SUCCESS);
14    }
```

Non-local goto in Signal Handler



- POSIX does not specify whether *longjmp()* will restore the signal context. If you want to save and restore **signal masks**, use *siglongjmp()*.
 - In a signal handler the signal that caused signal handler invocation is added to the signal mask. It will not be removed, if *longjmp()* is called.
- POSIX does not specify whether *setjmp()* will save the signal context. If you want to save signal masks, use *sigsetjmp()*.

```
2  #include <setjmp.h>
3  int sigsetjmp(sigjmp_buf env, int savemask);
4  /*Returns: 0 if called directly, nonzero if returning from a call to siglongjmp*/
5  void siglongjmp(sigjmp_buf env, int val);
```



Realtime Signals (R1: Ch22.8)

Realtime Signals Advantages



- Real time signals provide an increased range of signals that can be used for application-defined purposes.
- Realtime signals are queued.
- When sending a realtime signal, it is possible to specify data
- The order of delivery of different realtime signals is guaranteed.
 - If multiple different realtime signals are pending, then the lowest-numbered signal is delivered first.
- Realtime signals are not individually identified by different constants in the manner of standard signals.
 - a realtime signal number can be referred to by adding a value to SIGRTMIN ;
 - the expression $(\text{SIGRTMIN} + 1)$ refers to the second realtime signal

Send realtime signal



- The sending process sends the signal plus its accompanying data using the sigqueue() system call.

```
#define _POSIX_C_SOURCE 199309
#include <signal.h>
int sigqueue(pid_t pid, int sig, const union sigval value);
/*Returns 0 on success, or -1 on error*/
```

```
6 union sigval {
7     int sival_int;      /* Integer value for accompanying data */
8     void *sival_ptr;    /* Pointer value for accompanying data */
9 };
```

Receiving realtime signal



- The receiving process establishes a handler for the signal using a call to `sigaction()` that specifies the `SA_SIGINFO` flag.

```
1  struct sigaction act;  
2  sigemptyset(&act.sa_mask);  
3  act.sa_sigaction = handler;  
4  act.sa_flags = SA_RESTART | SA_SIGINFO;  
5  if (sigaction(SIGRTMIN + 5, &act, NULL) == -1)  
6      errExit("sigaction");
```

```
void handler(int sig, siginfo_t *siginfo, void *ucontext);
```

siginfo structure



```
1  typedef struct {
2      int      si_signo;      /* Signal number */
3      int      si_code;      /* Signal code */
4      int      si_trapno;    /* Trap number for hardware-generated signal
5                               (unused on most architectures) */
6      union sigval si_value; /* Accompanying data from sigqueue() */
7      pid_t    si_pid;      /* Process ID of sending process */
8      uid_t    si_uid;      /* Real user ID of sender */
9      int      si_errno;    /* Error number (generally unused) */
10     void     *si_addr;    /* Address that generated signal
11                             (hardware-generated signals only) */
12     int      si_overrun;   /* Overrun count (Linux 2.6, POSIX timers) */
13     int      si_timerid;   /* (Kernel-internal) Timer ID
14                             (Linux 2.6, POSIX timers) */
15     long     si_band;     /* Band event (SIGPOLL/SIGIO) */
16     int      si_fd;       /* File descriptor (SIGPOLL/SIGIO) */
17     int      si_status;    /* Exit status or signal (SIGCHLD) */
18     clock_t  si_utime;    /* User CPU time (SIGCHLD) */
19     clock_t  si_stime;    /* System CPU time (SIGCHLD) */
20 } siginfo_t;
```

- For real time signals, si_signo, si_code (source), si_value, si_pid, si_uid are set.



BITS Pilani
Pilani Campus



Thank You